

Once connected, follow these commands on both EC2 instances:

Switch to the root user

`sudo -i`

Update the instance

`yum update -y`

Install NGINX

`yum install nginx -y`

Start NGINX

`systemctl start nginx`

Check NGINX status

`systemctl status nginx`

Enable NGINX to start on boot

`systemctl enable nginx`

```
[root@ip-172-31-2-176 ~]# systemctl status nginx
● nginx.service - The nginx HTTP and reverse proxy server
   Loaded: loaded (/usr/lib/systemd/system/nginx.service; disabled; preset: disabled)
   Active: active (running) since Sat 2024-09-21 14:14:14 UTC; 4s ago
     Process: 11149 ExecStartPre=/usr/bin/rm -f /run/nginx.pid (code=exited, status=0/SUCCESS)
     Process: 11164 ExecStartPre=/usr/sbin/nginx -t (code=exited, status=0/SUCCESS)
     Process: 11213 ExecStart=/usr/sbin/nginx (code=exited, status=0/SUCCESS)
    Main PID: 11269 (nginx)
      Tasks: 2 (limit: 1112)
     Memory: 2.2M
        CPU: 56ms
    CGroup: /system.slice/nginx.service
            └─11269 "nginx: master process /usr/sbin/nginx"
              └─11279 "nginx: worker process"

Sep 21 14:14:13 ip-172-31-2-176.ap-south-1.compute.internal systemd[1]: Starting nginx.service - The nginx HTTP and reverse proxy server...
Sep 21 14:14:14 ip-172-31-2-176.ap-south-1.compute.internal nginx[11164]: nginx: the configuration file /etc/nginx/nginx.conf syntax is ok
Sep 21 14:14:14 ip-172-31-2-176.ap-south-1.compute.internal nginx[11164]: nginx: configuration file /etc/nginx/nginx.conf test is successful
Sep 21 14:14:14 ip-172-31-2-176.ap-south-1.compute.internal systemd[1]: Started nginx.service - The nginx HTTP and reverse proxy server.
[root@ip-172-31-2-176 ~]#
```

3. Configure Web Content for Each Server

Now, add unique content to each web server to differentiate between them.

On Web Server 1:

Go to the NGINX HTML directory

`cd /usr/share/nginx/html`

Replace the default index.html file with custom content

`echo "I am CloudDevOpshub batch36 student created Web Server 1" > index.html`

```
2. web_server_1 x 3. web_server_2 x +
[root@ip-172-31-2-176 ~]# cd /usr/share/nginx/html
[root@ip-172-31-2-176 html]# echo "I am CloudDevOpshub batch36 student created Web Server 1" > index.html
[root@ip-172-31-2-176 html]#
```

On Web Server 2:

Go to the NGINX HTML directory

`cd /usr/share/nginx/html`

Replace the default index.html file with custom content

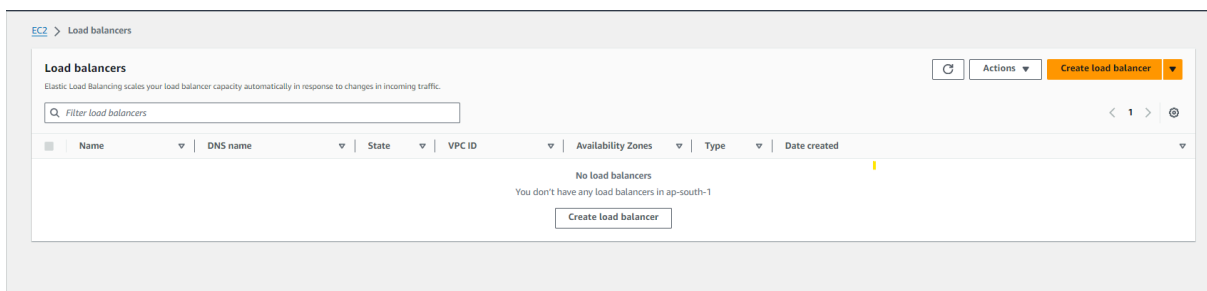
`echo "I am CloudDevOpshub batch36 student created Web Server 2" > index.html`

```
2. web_server_1 x 3. web_server_2 x +
[root@ip-172-31-11-134 ~]# cd /usr/share/nginx/html
[root@ip-172-31-11-134 html]# echo "I am CloudDevOpshub batch36 student created Web Server 2" > index.html
[root@ip-172-31-11-134 html]#
```

4. Create an Elastic Load Balancer (ELB)

Now, create an Elastic Load Balancer (ELB) that will distribute traffic between these two web servers.


Step 1: In the AWS Management Console, go to EC2 > Load Balancers and click Create Load Balancer.



Step 2: Select Application Load Balancer.

Load balancer types


Application Load Balancer [Info](#)



Choose an Application Load Balancer when you need a flexible feature set for your applications with HTTP and HTTPS traffic. Operating at the request level, Application Load Balancers provide advanced routing and visibility features targeted at application architectures, including microservices and containers.

Create

Network Load Balancer [Info](#)



Choose a Network Load Balancer when you need ultra-high performance, TLS offloading at scale, centralized certificate deployment, support for UDP, and static IP addresses for your applications. Operating at the connection level, Network Load Balancers are capable of handling millions of requests per second securely while maintaining ultra-low latencies.

Create

Gateway Load Balancer [Info](#)



Choose a Gateway Load Balancer when you need to deploy and manage a fleet of third-party virtual appliances that support GENEVE. These appliances enable you to improve security, compliance, and policy controls.

Create

► **Classic Load Balancer - previous generation**

Step 3: Configure the Load Balancer:

- Name: **MyELB**
- Scheme: Internet-facing
- Listeners: HTTP on port 80
- Availability Zones: Select the same region and Availability Zone where your EC2 instances are located.

Create Application Load Balancer [Info](#)

The Application Load Balancer distributes incoming HTTP and HTTPS traffic across multiple targets such as Amazon EC2 instances, microservices, and containers, based on request attributes. When the load balancer receives a connection request, it evaluates the listener rules in priority order to determine which rule to apply, and if applicable, it selects a target from the target group for the rule action.

► How Application Load Balancers work

Basic configuration

Load balancer name

Name must be unique within your AWS account and can't be changed after the load balancer is created.

MyELB

A maximum of 32 alphanumeric characters including hyphens are allowed, but the name must not begin or end with a hyphen.

Scheme [Info](#)

Scheme can't be changed after the load balancer is created.

☒ Internet-facing

An internet-facing load balancer routes requests from clients over the internet to targets. Requires a public subnet. [Learn more](#)

☐ Internal

An internal load balancer routes requests from clients to targets using private IP addresses. Compatible with the IPv4 and Dualstack IP address types.

Load balancer IP address type [Info](#)

Select the front-end IP address type to assign to the load balancer. The VPC and subnets mapped to this load balancer must include the selected IP address types. Public IPv4 addresses have an additional cost.

☒ IPv4

Includes only IPv4 addresses.

☐ Dualstack

Includes IPv4 and IPv6 addresses.

☐ Dualstack without public IPv4

Includes a public IPv6 address, and private IPv4 and IPv6 addresses. Compatible with internet-facing load balancers only.

Network mapping [Info](#)

The load balancer routes traffic to targets in the selected subnets, and in accordance with your IP address settings.

VPC [Info](#)

The load balancer will exist and scale within the selected VPC. The selected VPC is also where the load balancer targets must be hosted unless routing to Lambda or on-premises targets, or if using VPC peering. To confirm the VPC for your targets, view [target groups](#). For a new VPC, [create a VPC](#).

vpc-08b7f9d6407ed0f3c
IPv4 VPC CIDR: 172.31.0.0/16

Mappings [Info](#)

Select at least two Availability Zones and one subnet per zone. The load balancer routes traffic to targets in these Availability Zones only. Availability Zones that are not supported by the load balancer or the VPC are not available for selection.

Availability Zones

☐ ap-south-1a (aps1-az1)

☒ ap-south-1b (aps1-az3)

Subnet

subnet-0bcd66c14ab4374a
IPv4 subnet CIDR: 172.31.0.0/20

IPv4 address

Security groups [Info](#)

A security group is a set of firewall rules that control the traffic to your load balancer. Select an existing security group, or you can [create a new security group](#).

Security groups

Select up to 5 security groups

launch-wizard-3

sg-0d8f4d0b4ad8d3b52 VPC: vpc-08b7f9d6407ed0f3c

Listeners and routing [Info](#)

A listener is a process that checks for connection requests using the port and protocol you configure. The rules that you define for a listener determine how the load balancer routes requests to its registered targets.

▼ Listener HTTP:80

Remove

Protocol

HTTP

Port

80

1-65535

Default action

[Info](#)

Forward to

Select a target group

[Create target group](#)

Listener tags - *optional*

Consider adding tags to your listener. Tags enable you to categorize your AWS resources so you can more easily manage them.

Add listener tag

You can add up to 50 more tags.

Add listener

Step 4: Configure the target group:

- Target Type: Instance
- Protocol: HTTP
- Port: 80
- Health Checks: HTTP, path /

Listeners and routing [Info](#)

A listener is a process that checks for connection requests using the port and protocol you configure. The rules that you define for a listener determine how the load balancer routes requests to its registered targets.

▼ Listener HTTP:80

Remove

Protocol

HTTP

Port

80

1-65535

Default action

[Info](#)

Forward to

Select a target group

[Create target group](#)

⊗ You must select at least one target group.

Listener tags - *optional*

Consider adding tags to your listener. Tags enable you to categorize your AWS resources so you can more easily manage them.

Add listener tag

You can add up to 50 more tags.

Add listener

This is an optional step to create a target group. However, to ensure that your load balancer routes traffic to this target group you must register your targets.

0 pending

Cancel

Previous

Create target group

<input type="checkbox"/>	i-0157085c92d5721d7	Web_server_1	80	ap-south-1b	Unused	Target group is not co...	Septembe...	Normal
--------------------------	-------------------------------------	--------------	----	-------------	--------	---------------------------	-------------	--------


MyELB	MyELB-574212798.ap-sou...	Active	vpc-08b7f9d6407ed0f3c	2 Availability Zones	application	September 21, 2024, 19:59 (UTC+05:30)
---	---	--	-----------------------	----------------------	-------------	---------------------------------------

5. Test the Load Balancer

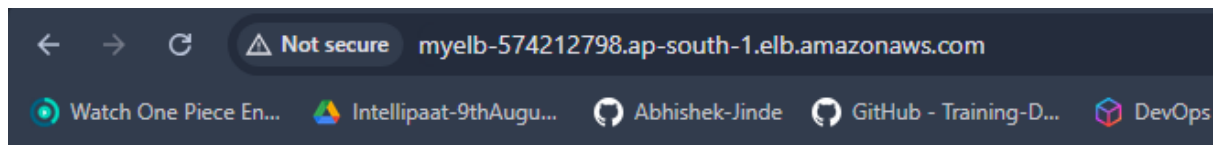
Once the Load Balancer is active, you can test it by visiting the Load Balancer's DNS name (found in the Load Balancer details) in your browser:

<http://<ELB-DNS-Name>>

You should see the content from either Web Server 1 or Web Server 2, and it will alternate as the load balancer distributes traffic between the two.



I am CloudDevOpshub batch36 student created Web Server 1



I am CloudDevOpshub batch36 student created Web Server 2